



# Data mesh

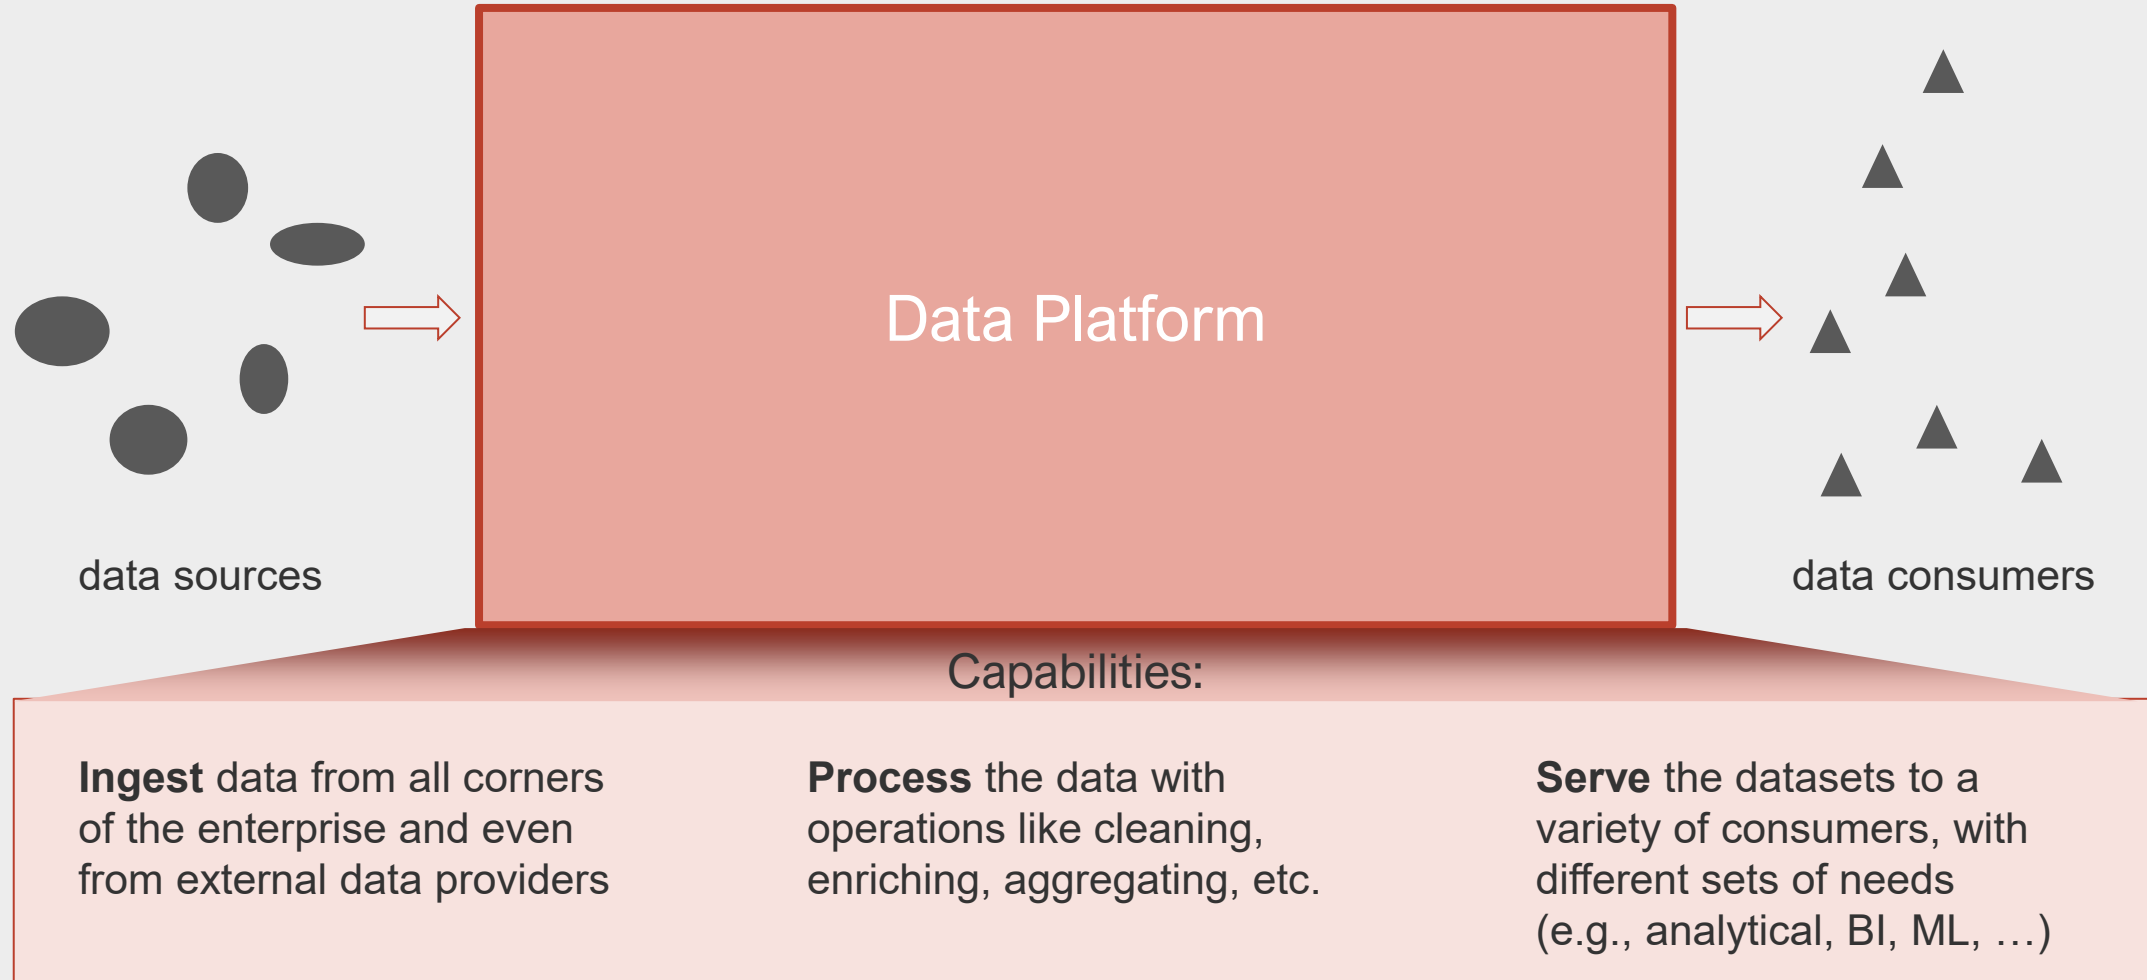
by Nicola Orecchini, 24/07/2025

# What you'll find in this presentation

Chapter	Content
1. Challenges of monolithic data architectures	Monolithic data architectures & their main points of failure
2. Data mesh: a new decentralized data architecture	A proposed evolution to solve some challenges


# 1. Challenges of monolithic data architectures

Most of the modern data architectures (such as data lake) are composed by a single, centralized, monolithic piece







There are 3  
main failure  
modes in  
such an  
architecture



Centralized and  
monolithic



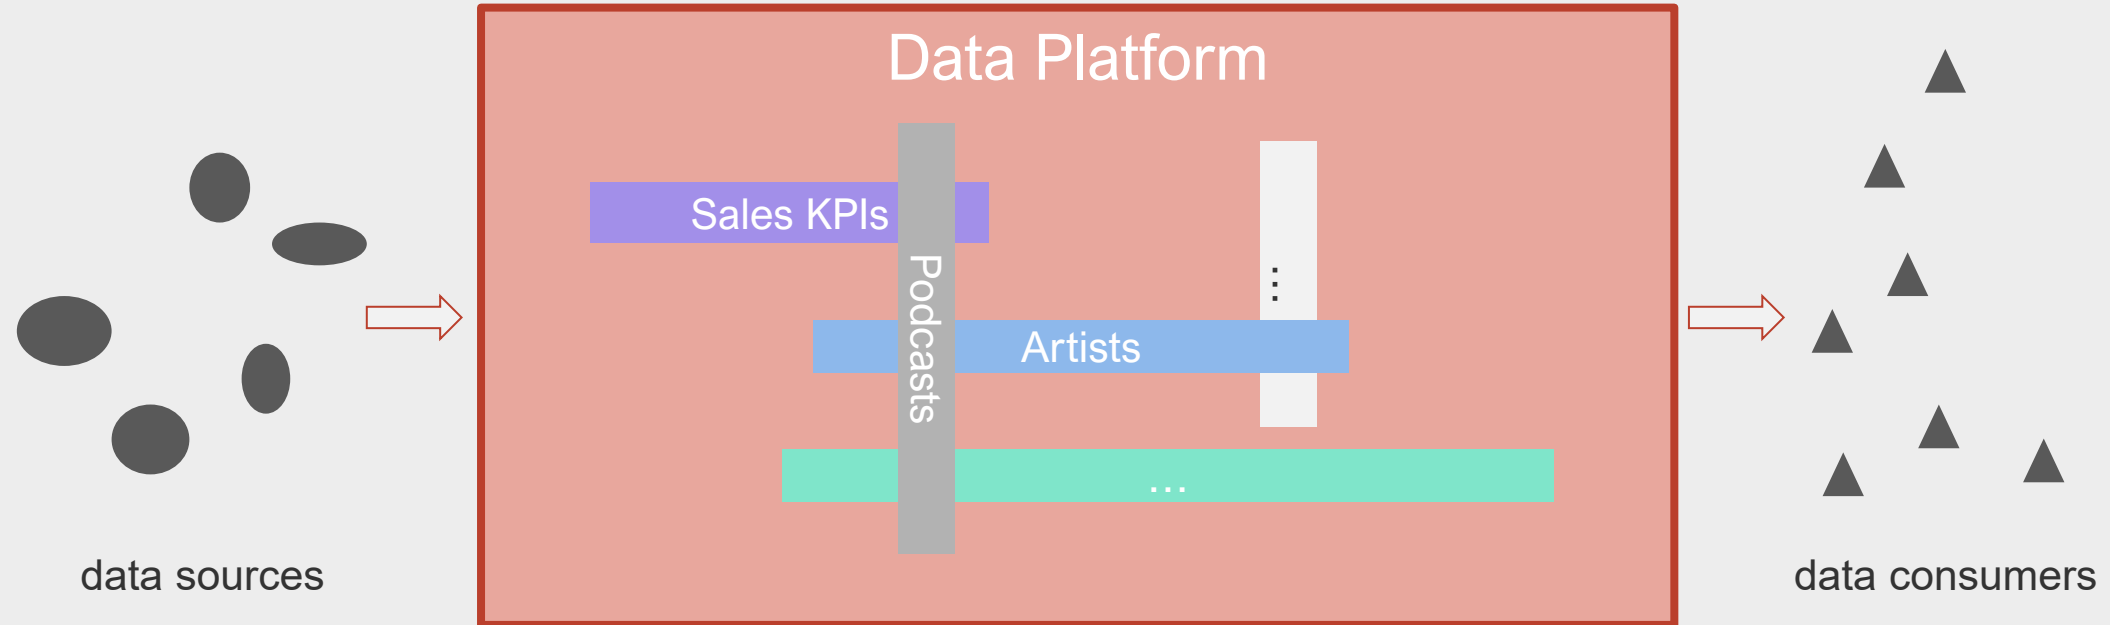
Coupled pipeline  
decomposition



Siloed and hyper-specialized  
ownership



# Inside the data platform, there are data from many disparate domains with no clear boundaries and ownership

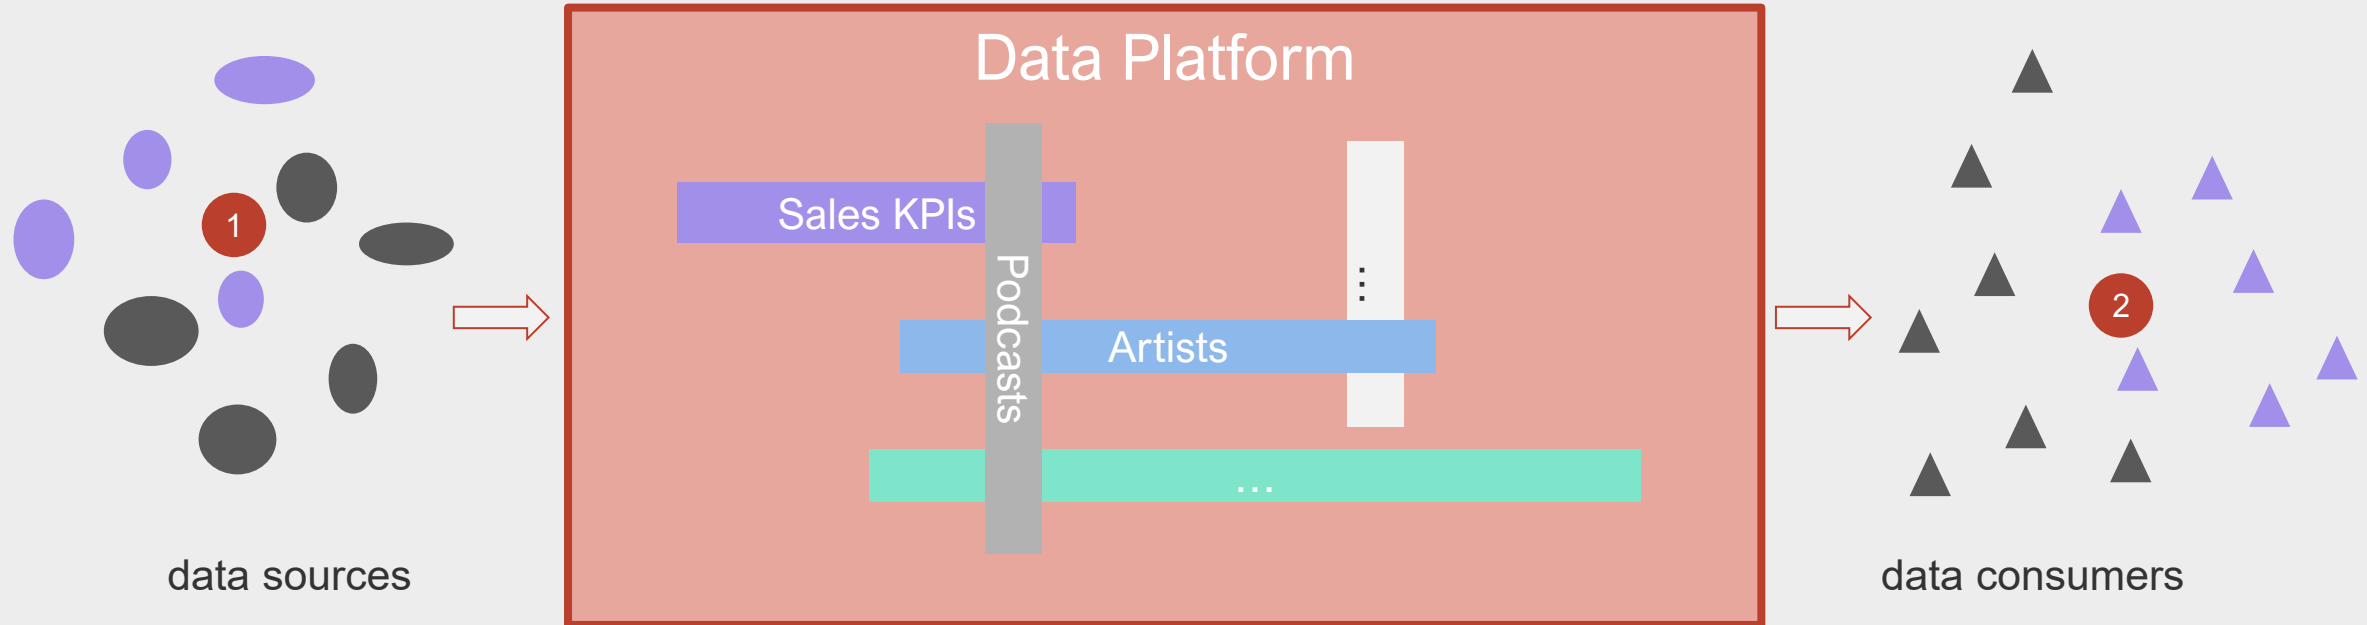


It is an accepted convention that the monolithic data platform hosts and owns **all the data in the organization**, even if that logically belong to different domains

Furthermore, **data from different domains are often in overlap**



## 2 pressure points threaten the efficiency of a centralized platform



### 1 Data source proliferation

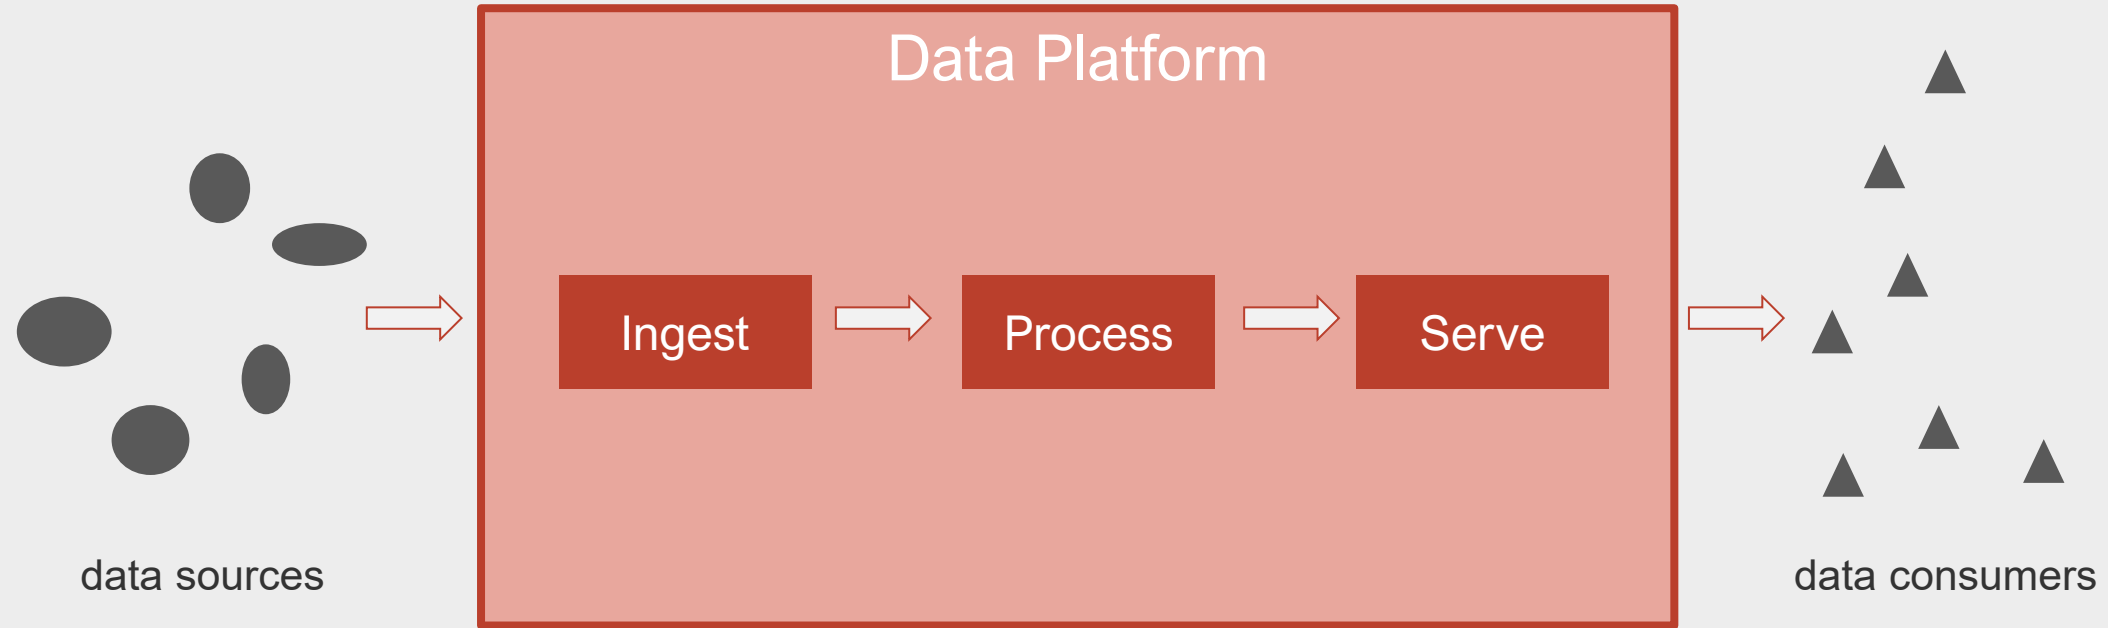
If the data sources start to become more and more, the assumption that we need to ingest and store all the data in one place will hinder our ability to timely & efficiently integrate the new data sources

### 2 Consumer proliferation

Organizations' need for rapid experimentation introduces a large number of use-cases for consumption of the data from the platform, which imply a growing number of transformations on the data, increasing the response time of the engineers 7



# If we zoom in on the architecture, we find the data platform is decomposed to a pipeline of data processing steps

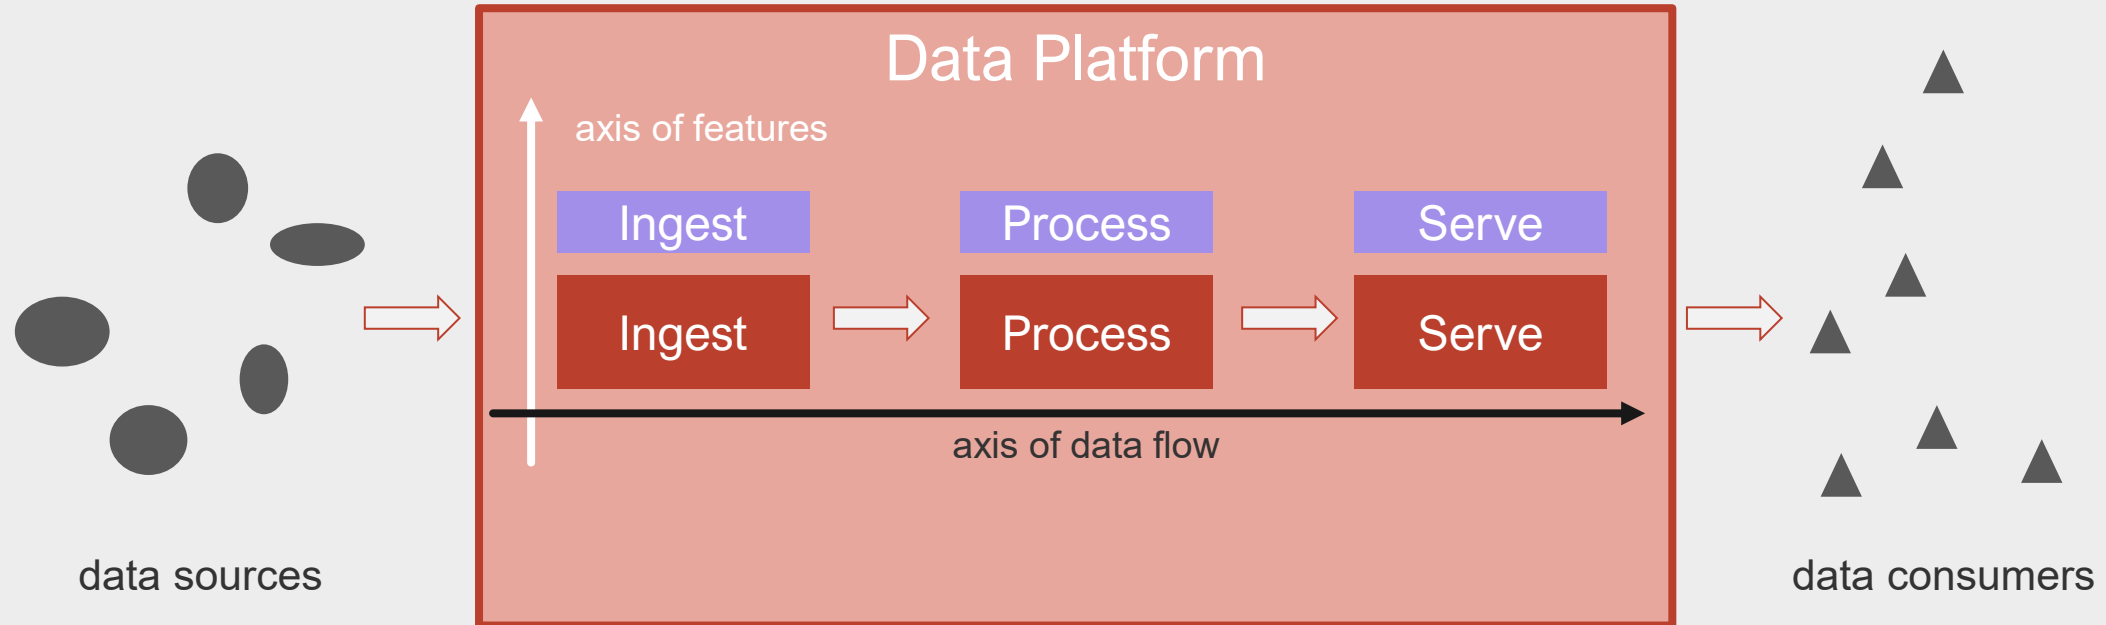


The grounding idea is that by decomposing the data platform architecture into separated capabilities, **teams can work parallelly** on different stages of the processing, and the work can **scale**





# However, the high level of coupling between stages slows the delivery of new features



On paper, with pipeline decomposition, it appears that we have achieved **modular components** that represent the **smallest units of change**

From the point of view of implementing a new feature, though, the **smallest unit of change** is still the **entire pipeline**: if you introduce a new feature, you must **change every component of the pipeline**. Thus, the pipeline decomposition is still **coupled**

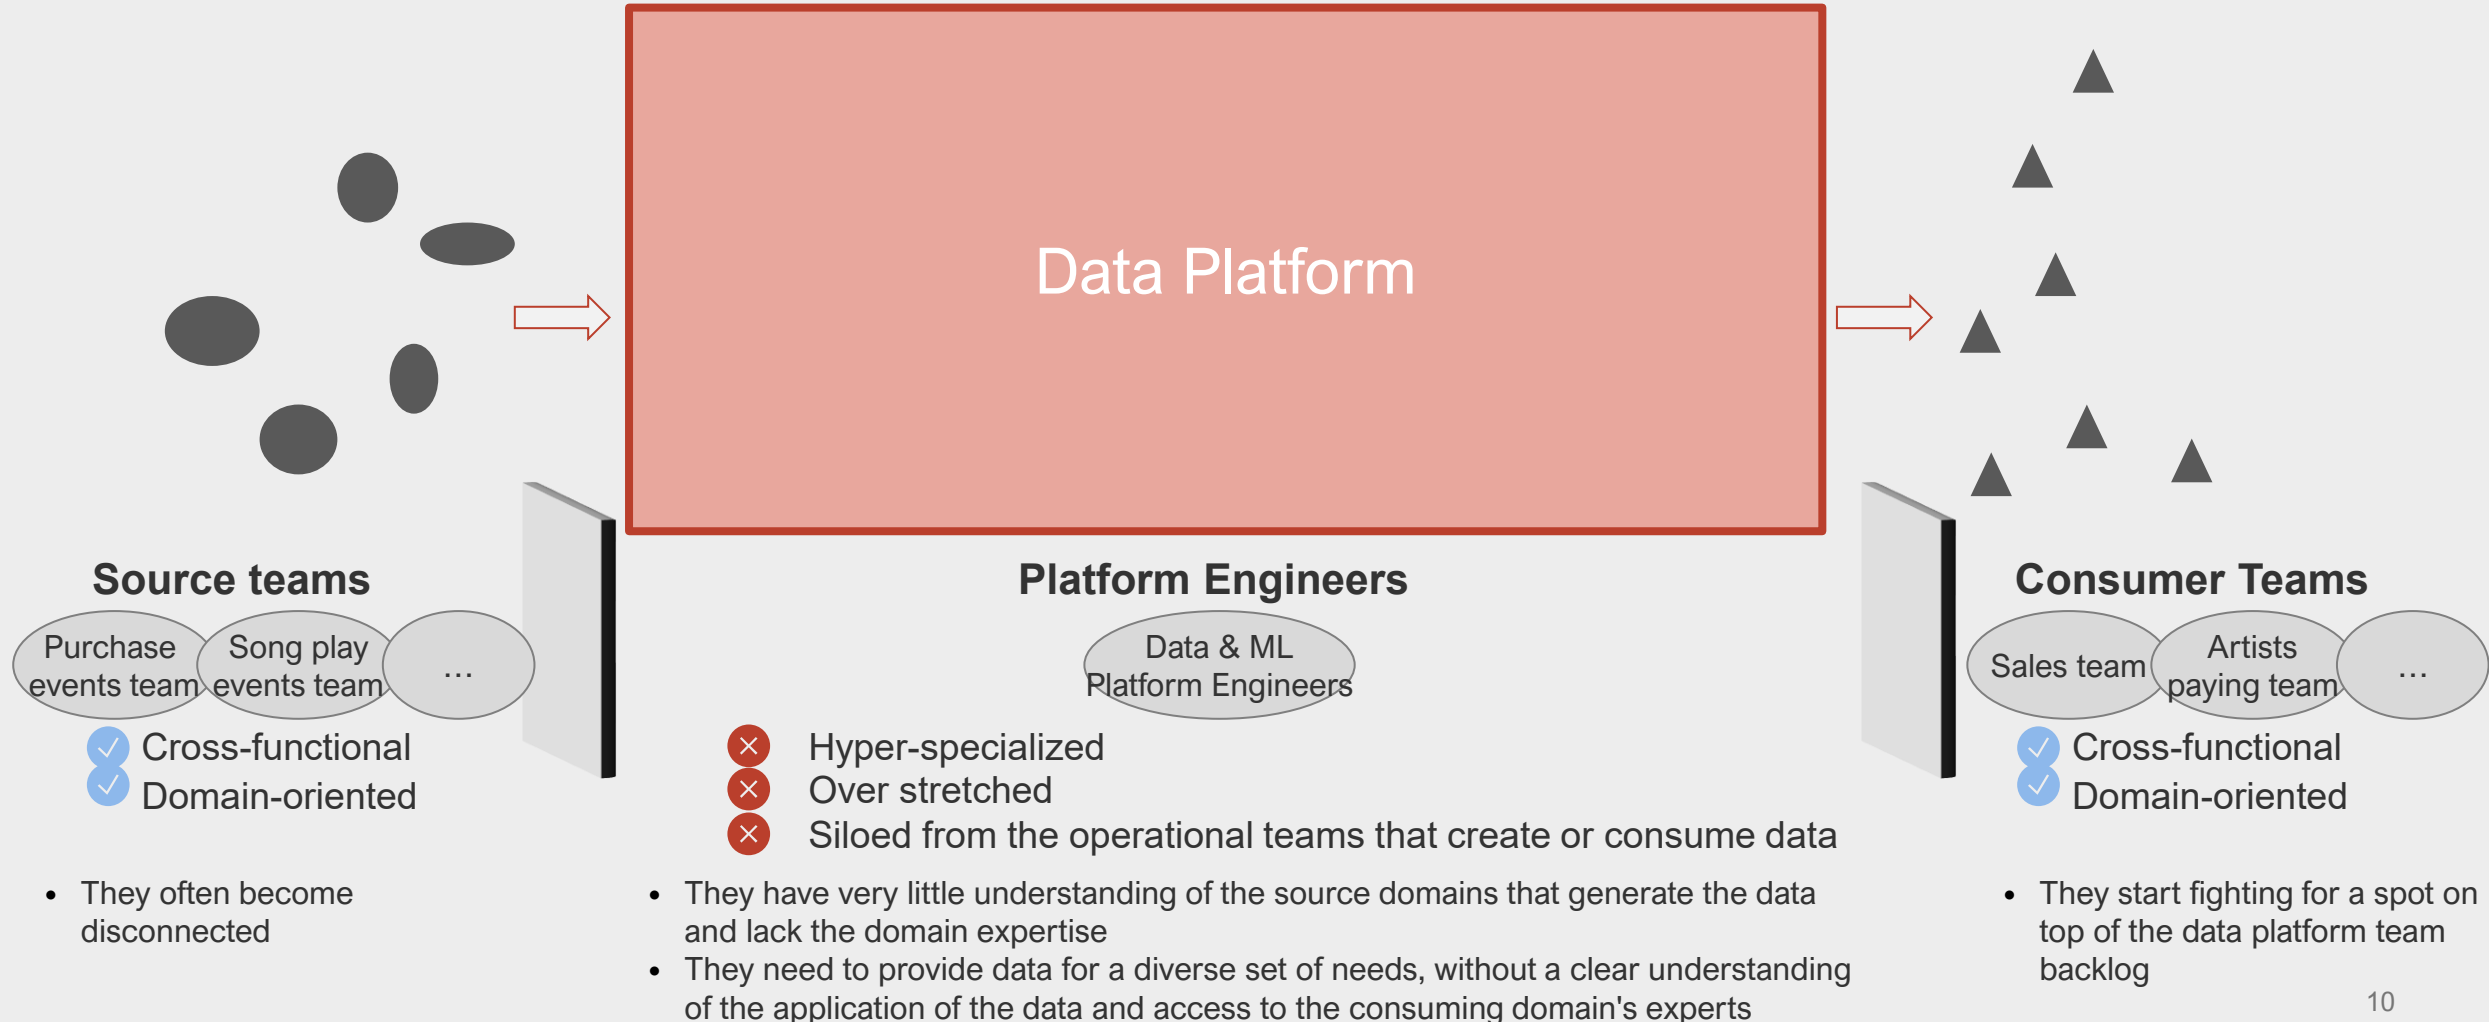
We could say that the pipeline decomposition acts on the **axis of data flow**, which is orthogonal to the **axis of features**

If you decompose the pipeline along the axis of data flow, you will achieve **benefits on parallelizing stages implementation**, but **no benefits on new features implementation**



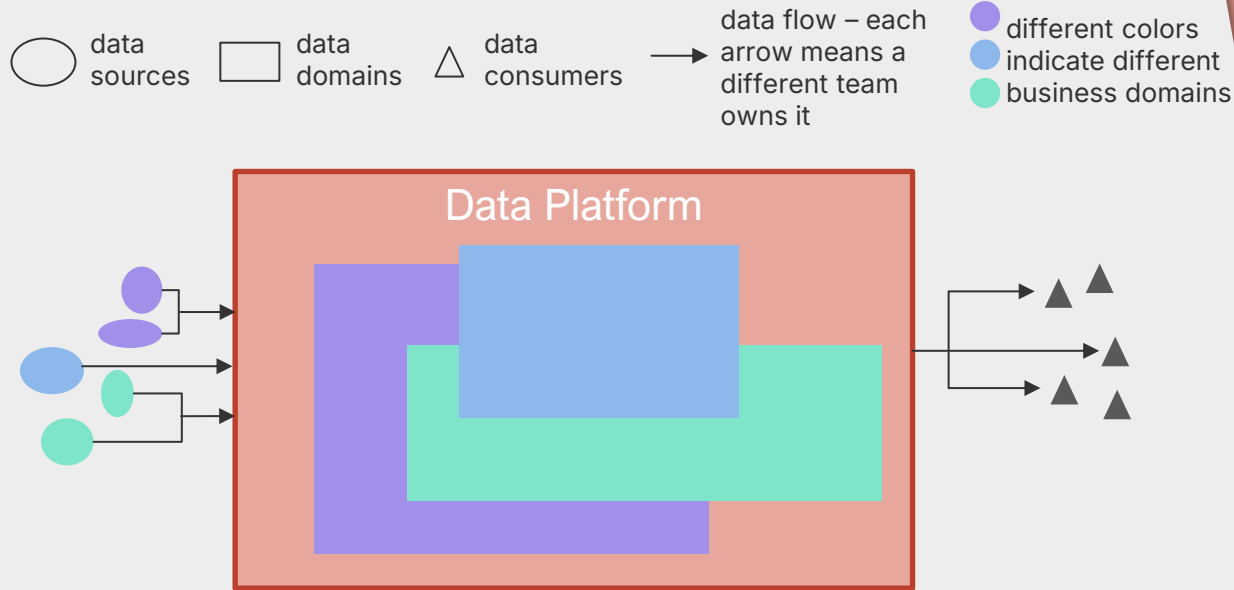
# Centralized platforms are built and owned by... guess who?

If you said *one single centralized team*, you're right!



## 2. Data mesh: a new decentralized data architecture

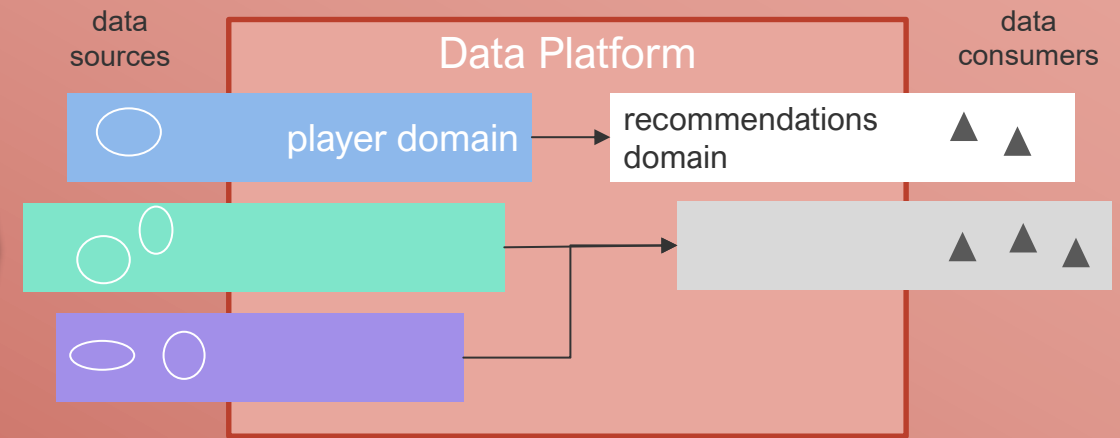
# To decentralize the monolithic data platform...



Push & ingest model

- each source operational system (managed by a domain team) feeds data organized by domains into the platform, not caring where these data will be consumed
- the centrally owned data platform ingests the data
- after ingestion, the concept of domain is lost, and one single platform team is responsible of providing data from the platform to consumers
- the architectural quantum is represented by a stage of the pipeline (e.g., ingest/process/serve)

# ...a paradigm shift is required to a domain-oriented data platform



Serve & pull model

- each domain owns, hosts and serves their datasets for access by any team downstream
- the physical location where the datasets actually reside and how they flow is a technical implementation of the *player* domain
  - the physical storage could still be centralized (e.g., Amazon S3 buckets), but domain datasets content and ownership remains within the domain generating them
- the architectural quantum is represented by a domain (e.g., player, recommendations, etc.)

# Disclaimer

This is my personal re-elaboration of Zhamak Dehghani's (the inventor of Data Mesh concept) article *How to Move Beyond a Monolithic Data Lake to a Distributed Data Mesh* (<https://martinfowler.com/articles/data-monolith-to-mesh.html>). I do not own the rights about the idea. For what concerns this deck, most of the sentences are copy-pasted from the article, and the graphics are also very similar to those of the article. But both the text and the graphics have been slightly re-elaborated by me, with the goal of making them more accessible and easier to understand.